

da24c026

September 14, 2024

#DA24C026 - Assignment 5

#Task 1

```
[1]: import numpy as np
import pandas as pd
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import OneHotEncoder, LabelEncoder
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
```

```
[3]: data = pd.read_csv("nursery.data", header = None)
```

```
[4]: data.columns = ["parents", "has_nursery", "form", "children", "housing", "
    ↪ "finance", "social", "health", "class"]
```

```
[5]: data.head()
```

```
[5]:  parents has_nursery    form children    housing    finance \
0   usual      proper  complete         1  convenient  convenient
1   usual      proper  complete         1  convenient  convenient
2   usual      proper  complete         1  convenient  convenient
3   usual      proper  complete         1  convenient  convenient
4   usual      proper  complete         1  convenient  convenient

      social    health    class
0   nonprob  recommended  recommend
1   nonprob    priority  priority
2   nonprob  not_recom  not_recom
3  slightly_prob  recommended  recommend
4  slightly_prob    priority  priority
```

```
[6]: data.describe()
```

```
[6]:
```

	parents	has_nursery	form	children	housing	finance \
count	12960	12960	12960	12960	12960	12960
unique	3	5	4	4	3	2
top	usual	proper	complete	1	convenient	convenient
freq	4320	2592	3240	3240	4320	6480

	social	health	class
count	12960	12960	12960
unique	3	3	5
top	nonprob	recommended	not_recom
freq	4320	4320	4320

```
[7]: data['class'].unique()
```

```
[7]: array(['recommend', 'priority', 'not_recom', 'very_recom', 'spec_prior'],
      dtype=object)
```

As we need a 3 class dataset we collapse recommend, very_recom and spec_prior class into recommend class

```
[8]: data["class"] = data["class"].replace("spec_prior", "recommend")
     data["class"] = data["class"].replace("very_recom", "recommend")
```

```
[9]: data['class'].describe()
```

```
[9]: count      12960
     unique        3
     top      recommend
     freq      4374
     Name: class, dtype: object
```

Train - Test Split

(80-20)

When Cross validation will be done using Grid Search, automatically a part of training data will be used as validation set.

```
[10]: def train_val_test(df):
      x_train, x_test, y_train, y_test = train_test_split(df.iloc[:, :-1], df.iloc[:,
      ↪, -1], test_size=0.2)
      return x_train, x_test, y_train, y_test
```

For model training and hyper parameter tuning, “Grid Search” is used.

Grid Search returns the combination of hyperparameters which yeild best accuracy.

##Decison Tree with categorical features

```
[11]: le = LabelEncoder()
x_encoded = data.copy()
for column in data.iloc[:, :-1]:
    x_encoded[column] = le.fit_transform(data[column])
```

```
[12]: x_encoded.head()
```

```
[12]:
```

	parents	has_nursery	form	children	housing	finance	social	health	\
0	2	3	0	0	0	0	0	2	
1	2	3	0	0	0	0	0	1	
2	2	3	0	0	0	0	0	0	
3	2	3	0	0	0	0	2	2	
4	2	3	0	0	0	0	2	1	

```

class
0 recommend
1 priority
2 not_recom
3 recommend
4 priority
```

```
[13]: x_train, x_test, y_train, y_test = train_val_test(x_encoded)
```

Grid Search implementation

```
[ ]: # DTREE Initialization
dtree = DecisionTreeClassifier()

# Hyperparameter grid
param_grid = {
    'max_depth': [None, 5, 10, 15, 20],
    'min_samples_split': [2, 5, 10, 15, 20],
    'min_samples_leaf': [1, 2, 4, 8, 16],
    'criterion': ['gini', 'entropy'] # Impurity measurement method
}

# Grid search for hyperparameter tuning
grid_search = GridSearchCV(estimator=dtree, param_grid=param_grid,
    scoring='accuracy', cv=5)
grid_search.fit(x_train, y_train)
```

```
[ ]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),
    param_grid={'criterion': ['gini', 'entropy'],
        'max_depth': [None, 5, 10, 15, 20],
        'min_samples_leaf': [1, 2, 4, 8, 16],
        'min_samples_split': [2, 5, 10, 15, 20]}},
    scoring='accuracy')
```

```
[ ]: # Best parameters from GridSearchCV
best_params = grid_search.best_params_
print("Best parameters found:", best_params)

# Final model with the best parameters
best_model = grid_search.best_estimator_

# Evaluation on the test set
y_pred = best_model.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)
print("Testing accuracy:", accuracy)
```

Best parameters found: {'criterion': 'entropy', 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2}
Testing accuracy: 0.9953703703703703

```
[ ]: dt_acc = [accuracy]
# we will repeat this process for 5 times to get mean accuracy and variance
```

```
[ ]: for i in range(4):
    x_train, x_test, y_train, y_test = train_val_test(x_encoded)
    grid_search = GridSearchCV(estimator=dtree, param_grid=param_grid,
    ↪scoring='accuracy', cv=5)
    grid_search.fit(x_train, y_train)
    best_params = grid_search.best_params_
    print("Best parameters found:", best_params)
    best_model = grid_search.best_estimator_
    y_pred = best_model.predict(x_test)
    accuracy = accuracy_score(y_test, y_pred)
    print("Testing accuracy:", accuracy)
    dt_acc.append(accuracy)
```

Best parameters found: {'criterion': 'gini', 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2}
Testing accuracy: 0.9934413580246914
Best parameters found: {'criterion': 'entropy', 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2}
Testing accuracy: 0.9949845679012346
Best parameters found: {'criterion': 'entropy', 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2}
Testing accuracy: 0.9942129629629629
Best parameters found: {'criterion': 'gini', 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2}
Testing accuracy: 0.9969135802469136

```
[ ]: mean_acc_dt = sum(dt_acc)/len(dt_acc)
variance_dt = sum((x - mean_acc_dt)**2 for x in dt_acc)/len(dt_acc)
```

```
print("Mean accuracy for Decison Tree:", mean_acc_dt)
print("Variance for Decison Tree:", variance_dt)
```

Mean accuracy for Decison Tree: 0.9949845679012345
Variance for Decison Tree: 1.3693606157597965e-06

```
[ ]: # Accuracy
acc = [100*i for i in dt_acc]
acc_dt = np.mean(acc)
print("Accuracy for Decison Tree:", acc_dt)
var_dt = np.var(acc)
print("Variance for Accuracy of Decison Tree:", var_dt)
```

Accuracy for Decison Tree: 99.49845679012346
Variance for Accuracy of Decison Tree: 0.013693606157597571

0.1 Decison Tree with One Hot Encoded features

One Hot Label Encoding - Converting categorical data to numerical

```
[ ]: one_hot_encoder = OneHotEncoder(sparse=False, drop='first') # Drop='first' to
    ↪avoid redundancy
x_oh_encoded = one_hot_encoder.fit_transform(data.iloc[:, :-1])
```

/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_encoders.py:975:
FutureWarning: `sparse` was renamed to `sparse_output` in version 1.2 and will
be removed in 1.4. `sparse_output` is ignored unless you leave `sparse` to its
default value.

```
warnings.warn(
```

```
[ ]: x_oh_encoded = pd.DataFrame(x_oh_encoded, columns=one_hot_encoder.
    ↪get_feature_names_out(data.columns[:-1]))
```

```
[ ]: x_oh_encoded["class"] = data["class"]
```

```
[ ]: x_oh_encoded.head()
```

```
[ ]:  parents_pretentious  parents_usual  has_nursery_improper  \
0                0.0            1.0                0.0
1                0.0            1.0                0.0
2                0.0            1.0                0.0
3                0.0            1.0                0.0
4                0.0            1.0                0.0

    has_nursery_less_proper  has_nursery_proper  has_nursery_very_crit  \
0                0.0            1.0                0.0
1                0.0            1.0                0.0
2                0.0            1.0                0.0
```

3	0.0	1.0	0.0
4	0.0	1.0	0.0

	form_completed	form_foster	form_incomplete	children_2	children_3 \
0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0	0.0

	children_more	housing_critical	housing_less_conv	finance_inconv \
0	0.0	0.0	0.0	0.0
1	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	0.0	0.0	0.0

	social_problematic	social_slightly_prob	health_priority \
0	0.0	0.0	0.0
1	0.0	0.0	1.0
2	0.0	0.0	0.0
3	0.0	1.0	0.0
4	0.0	1.0	1.0

	health_recommended	class
0	1.0	recommend
1	0.0	priority
2	0.0	not_recom
3	1.0	recommend
4	0.0	priority

Training and testing 5 Decision Tree models with OHE features. To get average accuracy and variance

```
[ ]: dt_oh_acc = []
for i in range(5):
    x_train, x_test, y_train, y_test = train_val_test(x_oh_encoded)
    dtree_oh = DecisionTreeClassifier()
    grid_search = GridSearchCV(estimator=dtree_oh, param_grid=param_grid,
    scoring='accuracy', cv=5)
    grid_search.fit(x_train, y_train)
    best_params = grid_search.best_params_
    print("Best parameters found:", best_params)
    best_model = grid_search.best_estimator_
    y_pred = best_model.predict(x_test)
    accuracy = accuracy_score(y_test, y_pred)
    print("Testing accuracy:", accuracy)
```

```
dt_oh_acc.append(accuracy)
```

Best parameters found: {'criterion': 'entropy', 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2}

Testing accuracy: 0.9922839506172839

Best parameters found: {'criterion': 'gini', 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2}

Testing accuracy: 0.9915123456790124

Best parameters found: {'criterion': 'entropy', 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2}

Testing accuracy: 0.9911265432098766

Best parameters found: {'criterion': 'entropy', 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2}

Testing accuracy: 0.9965277777777778

Best parameters found: {'criterion': 'entropy', 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2}

Testing accuracy: 0.9938271604938271

```
[ ]: mean_acc_dt_oh = sum(dt_oh_acc)/len(dt_oh_acc)
      variance_dt_oh = sum((x - mean_acc_dt_oh)**2 for x in dt_oh_acc)/len(dt_oh_acc)
      print("Mean accuracy for Decison Tree with OHE features:", mean_acc_dt_oh)
      print("Variance for Decison Tree with OHE features:", variance_dt_oh)
```

Mean accuracy for Decison Tree with OHE features: 0.9930555555555556

Variance for Decison Tree with OHE features: 3.869932174973301e-06

```
[ ]: # Accuracy
      acc = [100*i for i in dt_oh_acc]
      acc_dt_oh = np.mean(acc)
      print("Accuracy for Decison Tree OHE:", acc_dt_oh)
      var_dt_oh = np.var(acc)
      print("Variance for Accuracy of Decison Tree OHE:", var_dt_oh)
```

Accuracy for Decison Tree OHE: 99.30555555555554

Variance for Accuracy of Decison Tree OHE: 0.038699321749733465

0.2 KNN Implementation

```
[ ]: knn_acc = []
      param_grid = {
          'n_neighbors': [3, 5, 7, 9, 11, 13, 15],
          'weights': ['uniform', 'distance'], # Uniform - All neighbors are treated
          ↪equally, Distance - Neighbors have different influence depending upon their
          ↪distance from datapoint
          'p': [1, 2] # 1 for Manhattan distance, 2 for Euclidean distance
      }
```

```

for i in range(5):
    x_train, x_test, y_train, y_test = train_val_test(x_encoded)
    knn = KNeighborsClassifier()
    grid_search = GridSearchCV(estimator=knn, param_grid=param_grid,
    scoring='accuracy', cv=5)
    grid_search.fit(x_train, y_train)
    best_params = grid_search.best_params_
    print("Best parameters found:", best_params)
    best_model = grid_search.best_estimator_
    y_pred = best_model.predict(x_test)
    accuracy = accuracy_score(y_test, y_pred)
    print("Testing accuracy:", accuracy)
    knn_acc.append(accuracy)

```

```

Best parameters found: {'n_neighbors': 9, 'p': 1, 'weights': 'distance'}
Testing accuracy: 0.9629629629629629
Best parameters found: {'n_neighbors': 9, 'p': 1, 'weights': 'distance'}
Testing accuracy: 0.9564043209876543
Best parameters found: {'n_neighbors': 9, 'p': 1, 'weights': 'distance'}
Testing accuracy: 0.9645061728395061
Best parameters found: {'n_neighbors': 9, 'p': 1, 'weights': 'distance'}
Testing accuracy: 0.9560185185185185
Best parameters found: {'n_neighbors': 9, 'p': 1, 'weights': 'distance'}
Testing accuracy: 0.9560185185185185

```

```

[ ]: mean_acc_knn = sum(knn_acc)/len(knn_acc)
    variance_knn = sum((x - mean_acc_knn)**2 for x in knn_acc)/len(knn_acc)
    print("Mean accuracy for KNN:", mean_acc_knn)
    print("Variance for KNN:", variance_knn)

```

```

Mean accuracy for KNN: 0.9591820987654321
Variance for KNN: 1.4074645633287561e-05

```

```

[ ]: #Accuracy
    acc = [100*i for i in knn_acc]
    acc_knn = np.mean(acc)
    print("Accuracy for KNN:", acc_knn)
    var_knn = np.var(acc)
    print("Variance for Accuracy of KNN:", var_knn)

```

```

Accuracy for KNN: 95.91820987654322
Variance for Accuracy of KNN: 0.1407464563328742

```


0.3 Logistic Regression with L1 regularisation

```
[ ]: param_grid = {  
    'C': [0.001, 0.01, 0.1, 1, 10, 100], # regularization Coeff  
    'penalty': ['l1'], # L1 regularisation  
    'solver': ['liblinear'] # solver for L1 regularisation  
}
```

```
[ ]: lr_acc = []  
for i in range(5):  
    x_train, x_test, y_train, y_test = train_val_test(x_oh_encoded)  
    logreg = LogisticRegression()  
    grid_search = GridSearchCV(estimator=logreg, param_grid=param_grid,  
                               scoring='accuracy', cv=5)  
    grid_search.fit(x_train, y_train)  
    best_params = grid_search.best_params_  
    print("Best parameters found:", best_params)  
    best_model = grid_search.best_estimator_  
    y_pred = best_model.predict(x_test)  
    accuracy = accuracy_score(y_test, y_pred)  
    print("Testing accuracy:", accuracy)  
    lr_acc.append(accuracy)
```

```
Best parameters found: {'C': 100, 'penalty': 'l1', 'solver': 'liblinear'}  
Testing accuracy: 0.9104938271604939  
Best parameters found: {'C': 10, 'penalty': 'l1', 'solver': 'liblinear'}  
Testing accuracy: 0.9128086419753086  
Best parameters found: {'C': 10, 'penalty': 'l1', 'solver': 'liblinear'}  
Testing accuracy: 0.9089506172839507  
Best parameters found: {'C': 100, 'penalty': 'l1', 'solver': 'liblinear'}  
Testing accuracy: 0.9108796296296297  
Best parameters found: {'C': 10, 'penalty': 'l1', 'solver': 'liblinear'}  
Testing accuracy: 0.9112654320987654
```

```
[ ]: mean_acc_lr = sum(lr_acc)/len(lr_acc)  
variance_lr = sum((x - mean_acc_lr)**2 for x in lr_acc)/len(lr_acc)  
print("Mean accuracy for Logistic Regression:", mean_acc_lr)  
print("Variance for Logistic Regression:", variance_lr)
```

```
Mean accuracy for Logistic Regression: 0.9108796296296298  
Variance for Logistic Regression: 1.5479728699892934e-06
```

```
[ ]: #Accuracy  
acc = [100*i for i in lr_acc]  
acc_lr = np.mean(acc)  
print("Accuracy for Logistic Regression:", acc_lr)  
var_lr = np.var(acc)  
print("Variance for Accuracy of Logistic Regression:", var_lr)
```

Accuracy for Logistic Regression: 91.08796296296296
Variance for Accuracy of Logistic Regression: 0.015479728699892725

0.4 Visualizing average accuracies and variances for different models

```
[ ]: models = ['Logistic Regression', 'Neural Network Classification', 'Random_
↳Forest Classification',
              'Support Vector Classification', 'Xgboost Classification', 'Logistic_
↳Regression with L1', 'KNN', 'Decison Tree OHE', 'Decision Tree Categorical']

means = [92.253, 100, 98.025, 99.198, 99.969, acc_lr, acc_knn, acc_dt_oh,
↳acc_dt]

variances = [0.746, 0, 0.244, 0.095, 0.0038, var_lr, var_knn, var_dt_oh, var_dt]

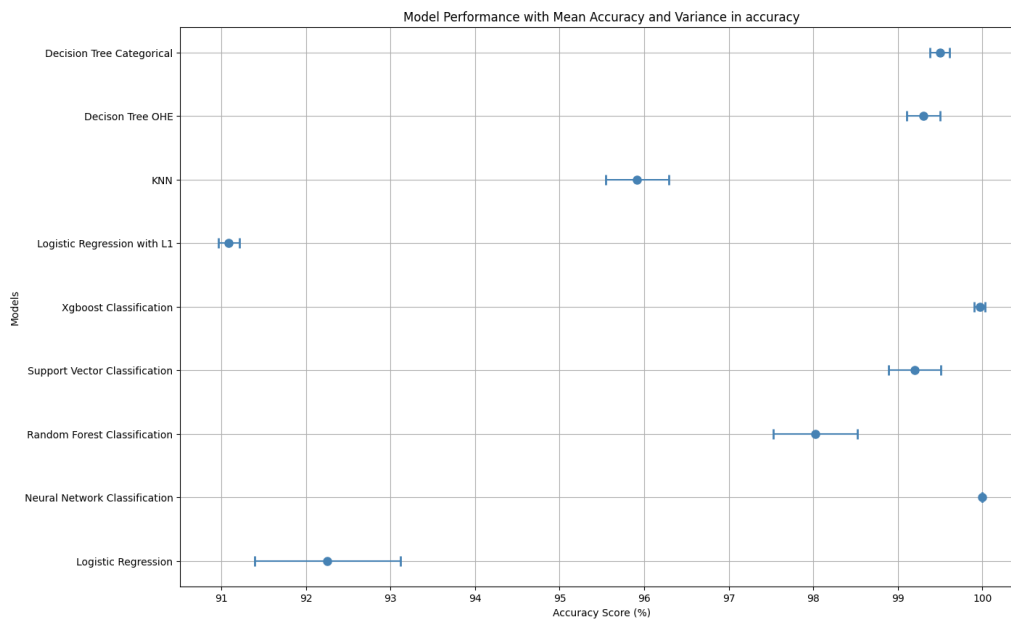
std_dev = np.sqrt(variances)

plt.figure(figsize=(15, 10))
plt.errorbar(means, models, xerr=std_dev, fmt='o', capsize=5, capthick=2,
↳markersize=8, color='steelblue')

plt.xlabel('Accuracy Score (%)')
plt.ylabel('Models')
plt.xticks(range(91,101))
plt.title('Model Performance with Mean Accuracy and Variance in accuracy')

plt.grid(True)

plt.show()
```



#Task 2

##Transforming Bipolar Sigmoid using unipolar sigmoid

Unipolar sigmoid is given by the following formula : - $\text{sigmoid}(x) = 1/(1+e^{-x})$. - It's value ranges from 0 to 1

Transformation from unipolar to bipolar is done by scaling unipolar sigmoid as follows: - $\text{bipolar_sigmoid}(x) = 2 * \text{sigmoid}(x) - 1$ - $\text{bipolar_sigmoid}(x) = -1 + 2/(1+e^{-x})$ - The value for this transformed bipolar sigmoid will range from -1 to 1. (Same as the range of Tanh function.)

```
[22]: def bipolar_sigmoid(x):
      z = 1/(1 + np.exp(-1*x))    #Sigmoid Function
      return 2 * z - 1
```

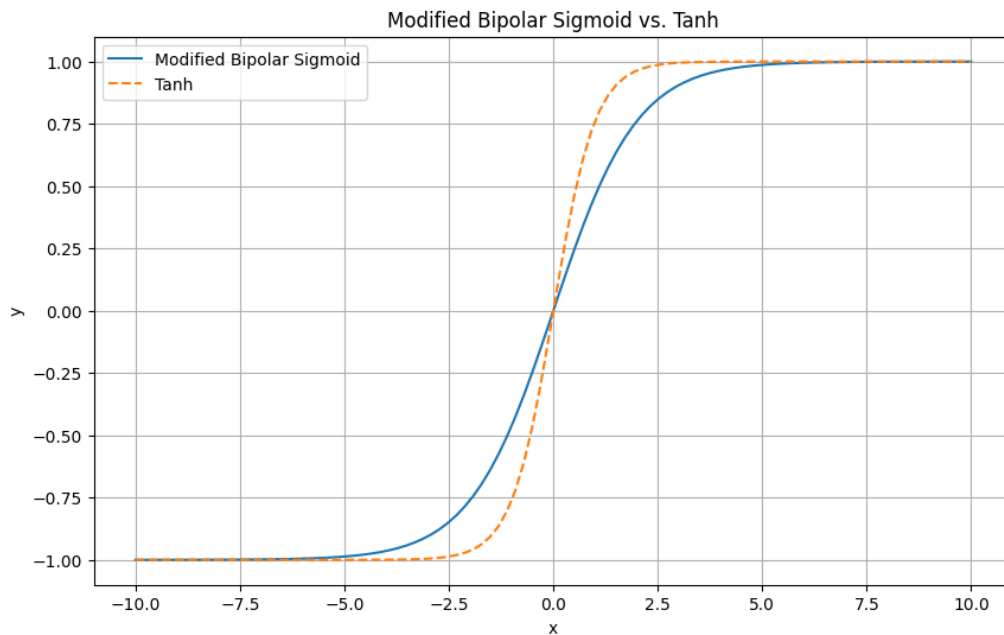
Tanh is another bipolar normalizer given by the following formula:

- $\tanh(x) = (e^x - e^{-x}) / (e^x + e^{-x})$

##Visualizing response for Modified Bipolar Sigmoid and Tanh

```
[30]: plt.figure(figsize=(10, 6))
      plt.plot(np.linspace(-10, 10, 100), bipolar_sigmoid(np.linspace(-10, 10, 100)),
               label='Modified Bipolar Sigmoid')
      plt.plot(np.linspace(-10, 10, 100), np.tanh(np.linspace(-10, 10, 100)),
               label='Tanh', linestyle="--")
      plt.xlabel('x')
      plt.ylabel('y')
```

```
plt.title('Modified Bipolar Sigmoid vs. Tanh')
plt.legend()
plt.grid(True)
plt.show()
```



The response for our Bipolar Sigmoid is very similar to Tanh normalizer, - They both range from -1 to 1 - Near origin, the steepness of slope for tanh is more than bipolar sigmoid

##Parameterizing and plotting Tanh and Bipolar sigmoid using different values of parameter “a”

```
[29]: a_values = [-5, -1, -0.1, -0.01, 0.001, 0.01, 0.1, 1, 5]
x = np.linspace(-10, 10, 100)

fig, ax = plt.subplots(3, 3, figsize=(15, 10))
ax = ax.ravel() # Flatten 2D plots

for idx, a in enumerate(a_values):
    bipolar_sigmoid_vals = bipolar_sigmoid(a * x)
    tanh_vals = np.tanh(a * x)

    ax[idx].plot(x, bipolar_sigmoid_vals, label='Bipolar Sigmoid', color='blue')
    ax[idx].plot(x, tanh_vals, label='Tanh', color='red', linestyle='--')

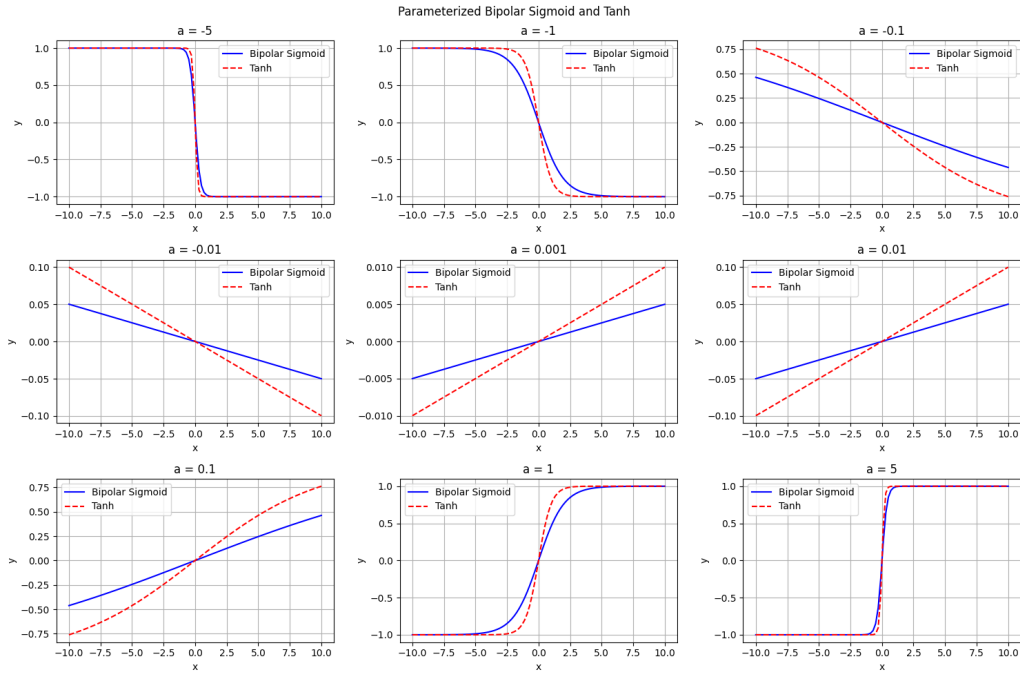
    ax[idx].set_xlabel('x')
    ax[idx].set_ylabel('y')
```

```

ax[idx].set_title(f'a = {a}')
ax[idx].legend()
ax[idx].grid(True)

plt.suptitle("Parameterized Bipolar Sigmoid and Tanh")
plt.tight_layout()
plt.show()

```



##Linearity Analysis from the graphs above

1. **When ($a = -5, 5$):** Both functions respond very sharply, with a steep transition near ($x = 0$). The non-linearity is significant.
2. **As (a) decreases ($a = 1, 0.1$):** The functions become less steep, and the transition becomes smoother, resulting in a wider range around ($x = 0$) where the output changes gradually.
3. **When (a) is very small ($a = -0.01$ to 0.01):** The curves almost appear linear across a broad range of (x), indicating a larger region of linear behavior.

0.4.1 Linear Range Analysis:

The linear range of the bipolar sigmoid increases as (a) tends closer to zero (e.g. -0.01 to 0.01). For very small (a) values, the bipolar sigmoid behaves almost **linearly over a wide range of (x)**, while for (a) values away from Zero (e.g. $-5, -1, 1, -5$), the function behaves more like a step function with a sharp transition at ($x = 0$).