

da24c026

October 13, 2024

DA24C026 - Assignment 7

TASK 1

```
[391]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import f1_score, make_scorer
from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.impute import SimpleImputer
from imblearn.over_sampling import SMOTE
from imblearn.over_sampling import RandomOverSampler
from imblearn.pipeline import Pipeline as ImbPipeline
import matplotlib.pyplot as plt
```

```
[305]: df = pd.read_csv("aps_failure_training_set.csv", skiprows = 20)
```

```
[306]: df.head()
```

```
[306]:   class  aa_000  ab_000      ac_000  ad_000  ae_000  af_000  ag_000  ag_001  ag_002  \
0   neg    76698    na  2130706438    280      0      0      0      0      0
1   neg    33058    na           0     na      0      0      0      0      0
2   neg    41040    na      228    100      0      0      0      0      0
3   neg        12     0       70     66      0     10      0      0      0
4   neg    60874    na     1368    458      0      0      0      0      0
```

```
   ...  ee_002  ee_003  ee_004  ee_005  ee_006  ee_007  ee_008  ee_009  ef_000  \
0   ...  1240520  493384  721044  469792  339156  157956   73224      0      0
1   ...   421400  178064  293306  245416  133654   81140   97576   1500      0
2   ...   277378  159812  423992  409564  320746  158022   95128    514      0
3   ...      240      46      58      44      10      0      0      0      4
4   ...   622012  229790  405298  347188  286954  311560  433954   1218      0
```

eg\_000

```
0      0
1      0
2      0
3     32
4      0
```

[5 rows x 171 columns]

```
[307]: df['class'].describe()
```

```
[307]: count      60000
unique         2
top           neg
freq         59000
Name: class, dtype: object
```

Data Preprocessing

```
[308]: df.replace('na', np.nan, inplace=True)
```

```
[309]: df.isna().sum()
```

```
[309]: class          0
aa_000            0
ab_000         46329
ac_000          3335
ad_000         14861
...
ee_007           671
ee_008           671
ee_009           671
ef_000          2724
eg_000          2723
Length: 171, dtype: int64
```

Dropping Columns which have more than 50% missing values as they don't add enough new info for training

```
[310]: df = df.dropna(axis=1, thresh=0.5*len(df))
```

```
[311]: df.shape
```

```
[311]: (60000, 163)
```

```
[312]: df.isna().sum()
```

```
[312]: class          0
      aa_000          0
      ac_000        3335
      ad_000       14861
      ae_000        2500
      ...
      ee_007         671
      ee_008         671
      ee_009         671
      ef_000        2724
      eg_000        2723
      Length: 163, dtype: int64
```

```
[313]: X = df.iloc[:,1:]
      y = df.iloc[:,0]
```

Handling missing values of remaining columns

```
[314]: imputer = SimpleImputer(strategy='median')
      X = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)
```

```
[315]: X.isna().sum()
```

```
[315]: aa_000      0
      ac_000      0
      ad_000      0
      ae_000      0
      af_000      0
      ..
      ee_007      0
      ee_008      0
      ee_009      0
      ef_000      0
      eg_000      0
      Length: 162, dtype: int64
```

Scaling data for faster convergence

```
[316]: std_scaler = StandardScaler()
      x_scaled = std_scaler.fit_transform(X)
```

Data dimensionality reduction

```
[317]: pca = PCA()
      X_pca = pca.fit_transform(x_scaled)

      explained_variance_ratio = pca.explained_variance_ratio_
```


```

cumulative_variance = np.cumsum(explained_variance_ratio)

plt.figure(figsize=(8, 5))
plt.plot(np.arange(1, len(cumulative_variance)+1), cumulative_variance,
        linestyle='--', color='b')
plt.title('Cumulative Explained Variance')
plt.xlabel('Number of Principal Components')
plt.ylabel('Cumulative Explained Variance')
plt.grid(True)
plt.show()

for i, cumulative_var in enumerate(cumulative_variance):
    print(f'Number of Components: {i+1}, Cumulative Explained Variance: {cumulative_var:.4f}')

```



```

Number of Components: 1, Cumulative Explained Variance: 0.3017
Number of Components: 2, Cumulative Explained Variance: 0.3512
Number of Components: 3, Cumulative Explained Variance: 0.3894
Number of Components: 4, Cumulative Explained Variance: 0.4199
Number of Components: 5, Cumulative Explained Variance: 0.4449
Number of Components: 6, Cumulative Explained Variance: 0.4653
Number of Components: 7, Cumulative Explained Variance: 0.4840
Number of Components: 8, Cumulative Explained Variance: 0.5024

```

Number of Components: 9, Cumulative Explained Variance: 0.5188  
 Number of Components: 10, Cumulative Explained Variance: 0.5338  
 Number of Components: 11, Cumulative Explained Variance: 0.5483  
 Number of Components: 12, Cumulative Explained Variance: 0.5612  
 Number of Components: 13, Cumulative Explained Variance: 0.5737  
 Number of Components: 14, Cumulative Explained Variance: 0.5854  
 Number of Components: 15, Cumulative Explained Variance: 0.5968  
 Number of Components: 16, Cumulative Explained Variance: 0.6079  
 Number of Components: 17, Cumulative Explained Variance: 0.6186  
 Number of Components: 18, Cumulative Explained Variance: 0.6291  
 Number of Components: 19, Cumulative Explained Variance: 0.6394  
 Number of Components: 20, Cumulative Explained Variance: 0.6495  
 Number of Components: 21, Cumulative Explained Variance: 0.6594  
 Number of Components: 22, Cumulative Explained Variance: 0.6689  
 Number of Components: 23, Cumulative Explained Variance: 0.6783  
 Number of Components: 24, Cumulative Explained Variance: 0.6872  
 Number of Components: 25, Cumulative Explained Variance: 0.6955  
 Number of Components: 26, Cumulative Explained Variance: 0.7036  
 Number of Components: 27, Cumulative Explained Variance: 0.7116  
 Number of Components: 28, Cumulative Explained Variance: 0.7192  
 Number of Components: 29, Cumulative Explained Variance: 0.7266  
 Number of Components: 30, Cumulative Explained Variance: 0.7338  
 Number of Components: 31, Cumulative Explained Variance: 0.7410  
 Number of Components: 32, Cumulative Explained Variance: 0.7480  
 Number of Components: 33, Cumulative Explained Variance: 0.7549  
 Number of Components: 34, Cumulative Explained Variance: 0.7617  
 Number of Components: 35, Cumulative Explained Variance: 0.7684  
 Number of Components: 36, Cumulative Explained Variance: 0.7751  
 Number of Components: 37, Cumulative Explained Variance: 0.7815  
 Number of Components: 38, Cumulative Explained Variance: 0.7878  
 Number of Components: 39, Cumulative Explained Variance: 0.7940  
 Number of Components: 40, Cumulative Explained Variance: 0.8002  
 Number of Components: 41, Cumulative Explained Variance: 0.8064  
 Number of Components: 42, Cumulative Explained Variance: 0.8123  
 Number of Components: 43, Cumulative Explained Variance: 0.8181  
 Number of Components: 44, Cumulative Explained Variance: 0.8240  
 Number of Components: 45, Cumulative Explained Variance: 0.8297  
 Number of Components: 46, Cumulative Explained Variance: 0.8352  
 Number of Components: 47, Cumulative Explained Variance: 0.8406  
 Number of Components: 48, Cumulative Explained Variance: 0.8457  
 Number of Components: 49, Cumulative Explained Variance: 0.8506  
 Number of Components: 50, Cumulative Explained Variance: 0.8554  
 Number of Components: 51, Cumulative Explained Variance: 0.8601  
 Number of Components: 52, Cumulative Explained Variance: 0.8646  
 Number of Components: 53, Cumulative Explained Variance: 0.8691  
 Number of Components: 54, Cumulative Explained Variance: 0.8734  
 Number of Components: 55, Cumulative Explained Variance: 0.8777  
 Number of Components: 56, Cumulative Explained Variance: 0.8820

Number of Components: 57, Cumulative Explained Variance: 0.8861  
Number of Components: 58, Cumulative Explained Variance: 0.8900  
Number of Components: 59, Cumulative Explained Variance: 0.8939  
Number of Components: 60, Cumulative Explained Variance: 0.8976  
Number of Components: 61, Cumulative Explained Variance: 0.9013  
Number of Components: 62, Cumulative Explained Variance: 0.9048  
Number of Components: 63, Cumulative Explained Variance: 0.9083  
Number of Components: 64, Cumulative Explained Variance: 0.9116  
Number of Components: 65, Cumulative Explained Variance: 0.9149  
Number of Components: 66, Cumulative Explained Variance: 0.9179  
Number of Components: 67, Cumulative Explained Variance: 0.9209  
Number of Components: 68, Cumulative Explained Variance: 0.9238  
Number of Components: 69, Cumulative Explained Variance: 0.9266  
Number of Components: 70, Cumulative Explained Variance: 0.9293  
Number of Components: 71, Cumulative Explained Variance: 0.9319  
Number of Components: 72, Cumulative Explained Variance: 0.9345  
Number of Components: 73, Cumulative Explained Variance: 0.9370  
Number of Components: 74, Cumulative Explained Variance: 0.9394  
Number of Components: 75, Cumulative Explained Variance: 0.9417  
Number of Components: 76, Cumulative Explained Variance: 0.9440  
Number of Components: 77, Cumulative Explained Variance: 0.9462  
Number of Components: 78, Cumulative Explained Variance: 0.9484  
Number of Components: 79, Cumulative Explained Variance: 0.9505  
Number of Components: 80, Cumulative Explained Variance: 0.9527  
Number of Components: 81, Cumulative Explained Variance: 0.9547  
Number of Components: 82, Cumulative Explained Variance: 0.9566  
Number of Components: 83, Cumulative Explained Variance: 0.9585  
Number of Components: 84, Cumulative Explained Variance: 0.9603  
Number of Components: 85, Cumulative Explained Variance: 0.9621  
Number of Components: 86, Cumulative Explained Variance: 0.9638  
Number of Components: 87, Cumulative Explained Variance: 0.9654  
Number of Components: 88, Cumulative Explained Variance: 0.9670  
Number of Components: 89, Cumulative Explained Variance: 0.9685  
Number of Components: 90, Cumulative Explained Variance: 0.9699  
Number of Components: 91, Cumulative Explained Variance: 0.9714  
Number of Components: 92, Cumulative Explained Variance: 0.9727  
Number of Components: 93, Cumulative Explained Variance: 0.9739  
Number of Components: 94, Cumulative Explained Variance: 0.9751  
Number of Components: 95, Cumulative Explained Variance: 0.9763  
Number of Components: 96, Cumulative Explained Variance: 0.9774  
Number of Components: 97, Cumulative Explained Variance: 0.9786  
Number of Components: 98, Cumulative Explained Variance: 0.9796  
Number of Components: 99, Cumulative Explained Variance: 0.9807  
Number of Components: 100, Cumulative Explained Variance: 0.9818  
Number of Components: 101, Cumulative Explained Variance: 0.9828  
Number of Components: 102, Cumulative Explained Variance: 0.9838  
Number of Components: 103, Cumulative Explained Variance: 0.9848  
Number of Components: 104, Cumulative Explained Variance: 0.9858

[illegible]

Number of Components: 153, Cumulative Explained Variance: 1.0000  
 Number of Components: 154, Cumulative Explained Variance: 1.0000  
 Number of Components: 155, Cumulative Explained Variance: 1.0000  
 Number of Components: 156, Cumulative Explained Variance: 1.0000  
 Number of Components: 157, Cumulative Explained Variance: 1.0000  
 Number of Components: 158, Cumulative Explained Variance: 1.0000  
 Number of Components: 159, Cumulative Explained Variance: 1.0000  
 Number of Components: 160, Cumulative Explained Variance: 1.0000  
 Number of Components: 161, Cumulative Explained Variance: 1.0000  
 Number of Components: 162, Cumulative Explained Variance: 1.0000

As it can be seen in the graph, if we have 60 components then more than 90% variance is captured

```
[318]: pca = PCA(n_components=60)
      X_pca = pca.fit_transform(x_scaled)
```

Splitting Data and training models

```
[319]: X_train, X_test, y_train, y_test = train_test_split(X_pca, y, test_size = 0.25,
      ↪ random_state = 42)
```

1.) SVC

```
[320]: svc = SVC()
      svc_param_grid = {
          'kernel': ['rbf', 'sigmoid'],
          'C': [0.1, 1, 10],
          'gamma': ['scale', 'auto']
      }
```

```
[321]: svc_grid = GridSearchCV(svc, svc_param_grid, cv=5, scoring = 'f1_macro')
      svc_grid.fit(X_train, y_train)
```

```
[321]: GridSearchCV(cv=5, estimator=SVC(),
      param_grid={'C': [0.1, 1, 10], 'gamma': ['scale', 'auto'],
      'kernel': ['rbf', 'sigmoid']},
      scoring='f1_macro')
```

```
[322]: svc_best = svc_grid.best_estimator_
      print ("Best parameters for SVC: ", svc_grid.best_params_)
```

Best parameters for SVC: {'C': 10, 'gamma': 'scale', 'kernel': 'rbf'}

```
[323]: y_train_pred_svc = svc_best.predict(X_train)
      y_test_pred_svc = svc_best.predict(X_test)

      f1_macro_train_svc = f1_score(y_train, y_train_pred_svc, average='macro')
      f1_macro_test_svc = f1_score(y_test, y_test_pred_svc, average='macro')
```



```
print("SVC Performance on Train Set:\n", f1_macro_train_svc)
print("SVC Performance on Test Set:\n", f1_macro_test_svc)
```

SVC Performance on Train Set:  
0.9790807359965954  
SVC Performance on Test Set:  
0.7851316110119919

## 2.) Logistic Regression

```
[324]: log_reg = LogisticRegression(solver='liblinear')
log_reg_param_grid = {
    'penalty': ['l1', 'l2'],
    'C': [0.1, 1, 10],
}
```

```
[325]: log_reg = GridSearchCV(log_reg, log_reg_param_grid, cv=5, scoring = "f1_macro")
log_reg.fit(X_train, y_train)
```

```
[325]: GridSearchCV(cv=5, estimator=LogisticRegression(solver='liblinear'),
    param_grid={'C': [0.1, 1, 10], 'penalty': ['l1', 'l2']},
    scoring='f1_macro')
```

```
[326]: log_reg_best = log_reg.best_estimator_
print ("Best parameters for Logistic Regression: ", log_reg.best_params_)
```

Best parameters for Logistic Regression: {'C': 10, 'penalty': 'l1'}

```
[327]: y_train_pred_log_reg = log_reg_best.predict(X_train)
y_test_pred_log_reg = log_reg_best.predict(X_test)

f1_macro_train_log_reg = f1_score(y_train, y_train_pred_log_reg,
    ↪average='macro')
f1_macro_test_log_reg = f1_score(y_test, y_test_pred_log_reg, average='macro')

print("Logistic Regression Performance on Train Set:\n", f1_macro_train_log_reg)
print("Logistic Regression Performance on Test Set:\n", f1_macro_test_log_reg)
```

Logistic Regression Performance on Train Set:  
0.8243212359917116  
Logistic Regression Performance on Test Set:  
0.835307388606881

## 3.) Decision Tree

```
[328]: dt = DecisionTreeClassifier()
dt_param_grid = {
    'max_depth': [10, 15, 20],
    'min_samples_leaf': [1, 2, 4]
```

```
}
```

```
[329]: dt_grid = GridSearchCV(dt, dt_param_grid, cv=5, scoring = "f1_macro")
dt_grid.fit(X_train, y_train)
```

```
[329]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),
                  param_grid={'max_depth': [10, 15, 20],
                              'min_samples_leaf': [1, 2, 4]},
                  scoring='f1_macro')
```

```
[330]: dt_best = dt_grid.best_estimator_
print ("Best parameters for Decision Tree: ", dt_grid.best_params_)
```

Best parameters for Decision Tree: {'max\_depth': 10, 'min\_samples\_leaf': 1}

```
[331]: y_train_pred_dt = dt_best.predict(X_train)
y_test_pred_dt = dt_best.predict(X_test)

f1_macro_train_dt = f1_score(y_train, y_train_pred_dt, average='macro')
f1_macro_test_dt = f1_score(y_test, y_test_pred_dt, average='macro')

print("Decision Tree Performance on Train Set:\n", f1_macro_train_dt)
print("Decision Tree Performance on Test Set:\n", f1_macro_test_dt)
```

Decision Tree Performance on Train Set:

0.9574379762054479

Decision Tree Performance on Test Set:

0.7885611433218309

Baseline Macro F1 Scores : SVC - 0.78, Logistic Reg - 0.83, Decision Tree - 0.79

## TASK 2

### Subtask 1 - Oversampling Minority Class

Using Smote for oversampling - Smote is used to oversample dataset based on no,of neighbors.

```
[332]: sm = SMOTE(random_state=12, sampling_strategy = 0.25)
x_train_res, y_train_res = sm.fit_resample(X_train,y_train)
```

```
[333]: y_train_res.describe()
```

```
[333]: count      55317
unique         2
top          neg
freq       44254
Name: class, dtype: object
```

Training models for different oversampling strategies

### 1.) SVC

```
[366]: clf = SVC(C=10, kernel='rbf', gamma='scale').fit(x_train_res,y_train_res)
```

```
[367]: y_train_pred_svc_os = clf.predict(X_train)
f1_macro_train_os_svc = f1_score(y_train, y_train_pred_svc_os, average='macro')
y_test_pred_svc_os = clf.predict(X_test)
f1_macro_test_os_svc = f1_score(y_test, y_test_pred_svc_os, average = "macro")
print(" Upon oversampling, F1 Score for SVC on training set :",\
      ↪f1_macro_train_os_svc)
print(" Upon oversampling, F1 Score for SVC on test set :",\
      ↪f1_macro_test_os_svc)
```

Upon oversampling, F1 Score for SVC on training set : 0.9790807359965954

Upon oversampling, F1 Score for SVC on test set : 0.7942521597522247

### 2.) Logistic Regression

```
[368]: sm = SMOTE(random_state=12, sampling_strategy = 0.06)
x_train_res, y_train_res = sm.fit_resample(X_train,y_train)
```

```
[369]: clf = LogisticRegression(solver='liblinear', C = 10, penalty='l1', max_iter =\
      ↪500).fit(x_train_res,y_train_res)
```

```
[370]: y_train_pred_logreg_os = clf.predict(X_train)
f1_macro_train_os_logreg = f1_score(y_train, y_train_pred_logreg_os,\
      ↪average='macro')
y_test_pred_logreg_os = clf.predict(X_test)
f1_macro_test_os_logreg = f1_score(y_test, y_test_pred_logreg_os, average =\
      ↪"macro")
print(" Upon oversampling, F1 Score for Logistic Regression on training set :",\
      ↪f1_macro_train_os_logreg)
print(" Upon oversampling, F1 Score for Logistic Regression on test set :",\
      ↪f1_macro_test_os_logreg)
```

Upon oversampling, F1 Score for Logistic Regression on training set :  
0.8488898122715383

Upon oversampling, F1 Score for Logistic Regression on test set :  
0.8400754227724834

### 3.) Decision Tree

```
[371]: sm = SMOTE(random_state=12, sampling_strategy = 0.02)
x_train_res, y_train_res = sm.fit_resample(X_train,y_train)
```

```
[372]: clf = DecisionTreeClassifier()
```

```
[373]: clf = DecisionTreeClassifier(max_depth=10, min_samples_leaf=1).fit(x_train_res,\
      ↪y_train_res)
```

```
[374]: y_train_pred_dtree_os = clf.predict(X_train)
f1_macro_train_os_dtree = f1_score(y_train, y_train_pred_dtree_os,
    ↪average='macro')
y_test_pred_dtree_os = clf.predict(X_test)
f1_macro_test_os_dtree = f1_score(y_test, y_test_pred_dtree_os, average =
    ↪"macro")
print(" Upon oversampling, F1 Score for Decision Tree on training set :",
    ↪f1_macro_train_os_dtree)
print(" Upon oversampling, F1 Score for Decision Tree on test set :",
    ↪f1_macro_test_os_dtree)
```

Upon oversampling, F1 Score for Decision Tree on training set :  
0.9626167990642506

Upon oversampling, F1 Score for Decision Tree on test set : 0.8039748989858151

Subtask 2 - Assigning class weights

SVC

```
[375]: clf = SVC(C=10, kernel='rbf', gamma='scale', class_weight= {'neg':1, 'pos':2}).
    ↪fit(X_train,y_train)
```

```
[376]: y_train_pred_svc_cw = clf.predict(X_train)
f1_macro_train_cw_svc = f1_score(y_train, y_train_pred_svc_cw, average='macro')
y_test_pred_svc_cw = clf.predict(X_test)
f1_macro_test_cw_svc = f1_score(y_test, y_test_pred_svc_cw, average = "macro")
print(" Upon assigning Class weights, F1 Score for SVC on training set :",
    ↪f1_macro_train_cw_svc)
print(" Upon assigning Class weights, F1 Score for SVC on test set :",
    ↪f1_macro_test_cw_svc)
```

Upon assigning Class weights, F1 Score for SVC on training set :  
0.9856597439968171

Upon assigning Class weights, F1 Score for SVC on test set : 0.7861070775130975

Logistic Regression

```
[377]: clf = LogisticRegression(solver='liblinear', C = 10, penalty='l1', max_iter =
    ↪500, class_weight={'neg':4, 'pos':11}).fit(X_train,y_train)
```

```
[378]: y_train_pred_logreg_cw = clf.predict(X_train)
f1_macro_train_cw_logreg = f1_score(y_train, y_train_pred_logreg_cw,
    ↪average='macro')
y_test_pred_logreg_cw = clf.predict(X_test)
f1_macro_test_cw_logreg = f1_score(y_test, y_test_pred_logreg_cw, average =
    ↪"macro")
print(" Upon assigning Class weights, F1 Score for Logistic Regression on
    ↪training set :", f1_macro_train_cw_logreg)
```

```
print(" Upon assigning Class weights, F1 Score for Logistic Regression on test_
↪set :", f1_macro_test_cw_logreg)
```

Upon assigning Class weights, F1 Score for Logistic Regression on training set :  
0.8457921283439009

Upon assigning Class weights, F1 Score for Logistic Regression on test set :  
0.8446141718090743

Decision Tree

```
[379]: clf = DecisionTreeClassifier(max_depth=10, min_samples_leaf=1,
↪class_weight={'neg':4, 'pos':11}).fit(X_train, y_train)
```

```
[410]: y_train_pred_dtree_cw = clf.predict(X_train)
f1_macro_train_cw_dtree = f1_score(y_train, y_train_pred_dtree_cw,
↪average='macro')
y_test_pred_dtree_cw = clf.predict(X_test)
f1_macro_test_cw_dtree = f1_score(y_test, y_test_pred_dtree_cw, average =
↪"macro")
print(" Upon assigning Class weights, F1 Score for Decision on training set :",
↪f1_macro_train_cw_dtree)
print(" Upon assigning Class weights, F1 Score for Decision Regression on test_
↪set :", f1_macro_test_cw_dtree)
```

Upon assigning Class weights, F1 Score for Decision on training set :  
0.9590298235571053

Upon assigning Class weights, F1 Score for Decision Regression on test set :  
0.816793893129771

Subtask - 3 Assigning sample weights instead of class weights, to penalize misclassifications

Main Difference between sample weights and class weights is Sample weights are assigned to each individual sample (data point) in the training dataset. Thus, it's parameter is specified while fitting the model. While, Class weights are assigned to each class (category) in a classification problem. Thus, it's parameter is specified while making object of class.

SVC

```
[381]: weights = {'neg': 0.475, 'pos': 1}
sw = np.array([weights[class_] for class_ in y_train])
```

```
[382]: clf = SVC(C=10, kernel='rbf', gamma='scale').fit(X_train,y_train,
↪sample_weight= sw)
```

```
[383]: y_train_pred_svc_sw = clf.predict(X_train)
f1_macro_train_sw_svc = f1_score(y_train, y_train_pred_svc_sw, average='macro')
y_test_pred_svc_sw = clf.predict(X_test)
f1_macro_test_sw_svc = f1_score(y_test, y_test_pred_svc_sw, average = "macro")
print(" Upon assigning Sample weights, F1 Score for SVC on training set :",
↪f1_macro_train_sw_svc)
```

```
print(" Upon assigning Sample weights, F1 Score for SVC on test set :",  
      ↪f1_macro_test_sw_svc)
```

Upon assigning Sample weights, F1 Score for SVC on training set :

0.9796171911702234

Upon assigning Sample weights, F1 Score for SVC on test set :

0.7865366157429332

Logistic Regression

```
[384]: weights = {'neg': 0.75, 'pos': 2.25}  
sw = np.array([weights[class_] for class_ in y_train])
```

```
[385]: clf = LogisticRegression(solver='liblinear', C = 10, penalty='l1', max_iter =  
      ↪500).fit(X_train,y_train, sample_weight=sw)
```

```
[386]: y_train_pred_logreg_sw = clf.predict(X_train)  
f1_macro_train_sw_logreg = f1_score(y_train, y_train_pred_logreg_sw,  
      ↪average='macro')  
y_test_pred_logreg_sw = clf.predict(X_test)  
f1_macro_test_sw_logreg = f1_score(y_test, y_test_pred_logreg_sw, average =  
      ↪"macro")  
print(" Upon assigning Sample weights, F1 Score for Logistic Regression on  
      ↪training set :", f1_macro_train_sw_logreg)  
print(" Upon assigning Sample weights, F1 Score for Logistic Regression on test  
      ↪set :", f1_macro_test_sw_logreg)
```

Upon assigning Sample weights, F1 Score for Logistic Regression on training set : 0.8431944395038975

Upon assigning Sample weights, F1 Score for Logistic Regression on test set : 0.8403600017355732

Decision Tree

```
[387]: weights = {'neg': 0.75, 'pos': 2.15}  
sw = np.array([weights[class_] for class_ in y_train])
```

```
[388]: clf = DecisionTreeClassifier(max_depth=10, min_samples_leaf=1).fit(X_train,  
      ↪y_train, sample_weight=sw)
```

```
[389]: y_train_pred_dtree_sw = clf.predict(X_train)  
f1_macro_train_sw_dtree = f1_score(y_train, y_train_pred_dtree_sw,  
      ↪average='macro')  
y_test_pred_dtree_sw = clf.predict(X_test)  
f1_macro_test_sw_dtree = f1_score(y_test, y_test_pred_dtree_sw, average =  
      ↪"macro")  
print(" Upon assigning Sample weights, F1 Score for Decision Tree on training  
      ↪set :", f1_macro_train_sw_dtree)
```

```
print(" Upon assigning Sample weights, F1 Score for Decision Tree on test set :  
↪", f1_macro_test_sw_dtree)
```

Upon assigning Sample weights, F1 Score for Decision Tree on training set :  
0.9590298235571053

Upon assigning Sample weights, F1 Score for Decision Tree on test set :  
0.816793893129771

Subtask 4 - A different way to handle Imbalance

Stratified K-fold cross validation which takes enough samples from both classes

```
[443]: skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)  
  
# Classifiers  
svc = SVC()  
logreg = LogisticRegression(solver='liblinear')  
dt = DecisionTreeClassifier()  
  
smote = SMOTE(sampling_strategy=0.06)  
  
svc_pipeline = ImbPipeline([('smote', smote), ('svc', svc)])  
logreg_pipeline = ImbPipeline([('smote', smote), ('logisticregression',  
↪logreg)])  
dt_pipeline = ImbPipeline([('smote', smote), ('decisiontreeclassifier', dt)])  
  
svc_pipeline.fit(X_train, y_train)  
logreg_pipeline.fit(X_train, y_train)  
dt_pipeline.fit(X_train, y_train)  
  
svc_test_pred = svc_pipeline.predict(X_test)  
logreg_test_pred = logreg_pipeline.predict(X_test)  
dt_test_pred = dt_pipeline.predict(X_test)  
  
svc_kvc = f1_score(y_test, svc_test_pred, average='macro')  
logreg_kcv = f1_score(y_test, logreg_test_pred, average='macro')  
dtree_kcv = f1_score(y_test, dt_test_pred, average='macro')  
  
print("SVC F1 Macro Score (Test):", svc_kvc)  
print("Logistic Regression F1 Macro Score (Test):", logreg_kcv)  
print("Decision Tree F1 Macro Score (Test):", dtree_kcv)
```

SVC F1 Macro Score (Test): 0.850736519594085

Logistic Regression F1 Macro Score (Test): 0.8394369441196272

Decision Tree F1 Macro Score (Test): 0.7771914047514632

```
[448]: x, width = np.arange(3), 0.1  
spacing = 0.01
```

```

plt.bar(x - 2*(width + spacing), [f1_macro_test_svc, f1_macro_test_log_reg,
    ↪f1_macro_test_dt], width, label='Baseline')
plt.bar(x - (width + spacing), [f1_macro_test_os_svc, f1_macro_test_os_logreg,
    ↪f1_macro_test_os_dtree], width, label='Oversampling')
plt.bar(x, [f1_macro_test_cw_svc, f1_macro_test_cw_logreg,
    ↪f1_macro_test_cw_dtree], width, label='Class weights')
plt.bar(x + (width + spacing), [f1_macro_test_sw_svc, f1_macro_test_sw_logreg,
    ↪f1_macro_test_sw_dtree], width, label='Sample weights')
plt.bar(x + 2*(width + spacing), [svc_kvc, logreg_kcv, dtree_kcv], width,
    ↪label='Stratified K Fold CV')

plt.ylabel('Macro F1 Score')
plt.title('Baseline model vs Imbalance handling techniques (Testing set)')
plt.xticks(x,['SVC', 'Logistic Regression', 'Decision Tree'])
plt.ylim(0.75, 0.87)
plt.legend()
plt.tight_layout()
plt.show()

```

